

Classification of poetry on the basis of different genre

In this project the poetry has been classified into different genres using two models namely 1) Naive Bayes classifier 2) Logistic Regression for classification

The data set has been taken from Kaggle (<https://www.kaggle.com/ultrajack/modern-renaissance-poetry>)

In []:

```
import nltk
from nltk.corpus import stopwords
set(stopwords.words('english'))
```

In [115]:

```
import numpy as np
import sklearn as sk
import pandas as pd
dataset=pd.read_csv("C:/Users/Nirvan Dogra/Desktop/poems-from-poetryfoundation-org/all.csv")
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split
print(dataset.describe())
X_train, X_test, y_train, y_test = train_test_split(dataset['content'], dataset['type'],
                                                    test_size=0.2, random_state=42) #splitting the data into test and training
```

	author	content
\		
count	573	573
unique	67	506
top	WILLIAM SHAKESPEARE	Originally published in Poetry, March 1914.
freq	71	4

	poem name	age	type
count	571	573	573
unique	508	2	3
top	Canto IV	Renaissance	Love
freq	3	315	326

About Data

- > The data set contains 573 data values -> It has 5 columns namely (author, content, poem name, age, type)
- > The most common value is William Shakespeare containing a total of 71 data values

Pre-pocessing of data

For the pre-porcessing of data a library named nlkt has been used. For the pre-processing of data the following steps have been take: 1) Conversion to lower case 2) Bag of words 3) Removal of stop words 4) Removal of punctuation 5) Lemmatizer 6) conversion of text to sparse matrix

In [116]:

```

from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.feature_extraction.text import CountVectorizer

def text_process(example_sent):
    example_sent = example_sent.lower()
    stop_words = set(stopwords.words('english'))

    word_tokens = word_tokenize(example_sent)

    punctuation=['.',',','!',';','(',')','-', '_']
    filtered_sentence = []
    for w in word_tokens:
        if not (w in stop_words) and not (w in punctuation):
            filtered_sentence.append(w)
    return(filtered_sentence)

#     Lemmatizer = WordNetLemmatizer()
#     base_form=[];
#     for w in filtered_sentence:
#         print(w," ", Lemmatizer.Lemmatize(w, pos='a'))
#         base_form.append(Lemmatizer.Lemmatize(w, pos='a'))
#     print(base_form);

# example_sent = "This is a sample sentence, showing off the stop words filtration."
# for example_sent in X_train:

#     example_sent = example_sent.lower()
#     stop_words = set(stopwords.words('english'))

#     word_tokens = word_tokenize(example_sent)
#     print(word_tokens)
#     punctuation=['.',',','!',';','(',')','-', '_']
#     filtered_sentence = []
#     for w in word_tokens:
#         if not (w in stop_words) and not (w in punctuation):
#             filtered_sentence.append(w)
#     print(filtered_sentence)

#     Lemmatizer = WordNetLemmatizer()
#     base_form=[];
#     for w in filtered_sentence:
#         print(w," ", Lemmatizer.Lemmatize(w, pos='a'))
#         base_form.append(Lemmatizer.Lemmatize(w, pos='a'))
#     print(base_form);

#     word2count={}
#     for w in base_form:
#         if w not in word2count.keys():
#             word2count[w] = 1
#         else:
#             word2count[w] += 1

#     print(word2count);

```

```
#           # Importing necessary libraries

#           # instantiating the model with Multinomial Naive Bayes..
#           model = MultinomialNB()
#           # training the model...
#           model = model.fit(base_form, y_train)

# filtered_sentence = [w for w in word_tokens if not w in stop_words]

#

# for w in word_tokens:
#     if w not in stop_words:
#         filtered_sentence.append(w)

# print(word_tokens)
# print(filtered_sentence)
```

In [117]:

```
bow_transformer=CountVectorizer(analyzer=text_process).fit(X_train)
# transforming into Bag-of-Words and hence textual data to numeric..
text_bow_train=bow_transformer.transform(X_train)
# transforming into Bag-of-Words and hence textual data to numeric..
text_bow_test=bow_transformer.transform(X_test)
```

Naive Bayes

In [111]:

```
from sklearn.naive_bayes import MultinomialNB
# instantiating the model with Multinomial Naive Bayes..
model = MultinomialNB()
# training the model...
model = model.fit(text_bow_train, y_train)
```

In [73]:

```
model.score(text_bow_train, y_train)
```

Out[73]:

0.868995633187773

In [74]:

```
# Importing necessary libraries
from sklearn.metrics import classification_report

# getting the predictions of the Validation Set...
predictions = model.predict(text_bow_test)
# getting the Precision, Recall, F1-Score
print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
Love	0.72	0.87	0.79	68
Mythology & Folklore	0.20	0.11	0.14	9
Nature	0.68	0.50	0.58	38
micro avg	0.69	0.69	0.69	115
macro avg	0.53	0.49	0.50	115
weighted avg	0.67	0.69	0.67	115

Inferring form the results

The model has an accuracy of 86.89% The model works best at classifying 'Love' and worst for the classification of 'Mythology & Folklore'. To compensate for this, we can use weighted metrics on the Mythology and Folklore classes.

Logistic regression

The preprocessing of the data remains the same.

In [113]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
lm=LogisticRegression(solver='lbfgs', multi_class='multinomial').fit(text_bow_train, y_train)
predicted_classes = lm.predict(text_bow_test)
#print(predicted_classes)
accuracy = accuracy_score(y_test, predicted_classes)
parameters = model.coef_
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.p
y:758: ConvergenceWarning: lbfgs failed to converge. Increase the number o
f iterations.
"of iterations.", ConvergenceWarning)
```

Explanation of non convergence:

The model doesnt converge due to the presence of mulitple local minima, each sub-localized to the genre. Since there are multiple genre to be covered, it only makes sense that the minima is different for each genre.

In [114]:

```
print(accuracy)
```

```
0.6782608695652174
```

Inferring from the results

The accuracy of the model is 67.82% This model performs worse than the previous one.

Final comments

The accuracy of the above model can be improved with the use of explicit minima, i.e. using different minima in training models for each genre. The accuracy of the earlier model can be improved with a change in either pruning via AdaBoost, or by using weighted metrics and post pruning the model.